Howard Yen

Website: howard-yen.github.io
Email: hyen@cs.princeton.edu
Google Scholar: Howard Yen
GitHub: github.com/howard-yen

Education

Princeton University — Princeton, NJ
Ph.D. in Computer Science, Advisor: Danqi Chen — 2024–Current

Princeton University — Princeton, NJ
M.S.E. in Computer Science, GPA: 4.00/4.00 — 2023–2024
Thesis: "Long-Context Language Modeling with Parallel Context Encoding", Advised by Danqi Chen.

Princeton University — Princeton, NJ
B.S.E. in Computer Science, Highest Honors (*summa cum laude*), GPA: 3.99/4.00 — 2019–2023
Thesis: "How to Answer a Question? Rethinking Open-Domain Question Answering with Multi-Type Questions", Advised by Danqi Chen.

Publications

1. **Howard Yen**, Tianyu Gao, Minmin Hou, Ke Ding, Daniel Fleischer, Peter Izsak, Moshe Wasserblat, and Danqi Chen. "HELMET: How to Evaluate Long-Context Language Models Effectively and Thoroughly". To appear in *Proceedings of The International Conference on Learning Representations (ICLR)*, 2025. [Paper] [Code] [Site]

2. Hongjin Su*, **Howard Yen***, Mengzhou Xia*, Weijia Shi, Niklas Muennighoff, Han-yu Wang, Haisu Liu, Quan Shi, Zachary S. Siegel, Michael Tang, Ruoxi Sun, Jinsung Yoon, Sercan O. Arik, Danqi Chen, Tao Yu. "BRIGHT: A Realistic and Challenging Benchmark for Reasoning-Intensive Retrieval". To appear in *Proceedings of The International Conference on Learning Representations (ICLR)*, 2025. **Spotlight, top 5.1%.** [Paper] [Code] [Site]

3. Kenneth Enevoldsen, ..., **Howard Yen**, ..., Niklas Muennighoff. "MMTEB: Massive Multilingual Text Embedding Benchmark". To appear in *Proceedings of The International Conference on Learning Representations (ICLR)*, 2025. [Paper] [Code]

4. Xi Ye, Fangcong Yin, Yinghui He, Joie Zhang, **Howard Yen**, Tianyu Gao, Greg Durrett, and Danqi Chen. "LongProc: Benchmarking Long-Context Language Models on Long Procedural Generation". *Preprint*, 2025. [Paper] [Code] [Site]

5. Tianyu Gao*, Alexander Wettig*, **Howard Yen**, and Danqi Chen. "How to Train Long-Context Language Models (Effectively)". *Preprint*, 2024. [Paper] [Code]

6. **Howard Yen**, Tianyu Gao, and Danqi Chen. "Long-Context Language Modeling with Parallel Context Encoding". In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024. [Paper] [Code]

7. Ryan Liu, **Howard Yen**, Raja Marjieh, Thomas L. Griffiths, and Ranjay Krishna. "Optimizing Interpersonal Communication by Simulating Audiences with Large Language Models". *Preprint*, 2023. [Paper] [Code]

8. Tianyu Gao, **Howard Yen**, Jiatong Yu, and Danqi Chen. "Enabling Large Language Models to Generate Text with Citations". In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023. [Paper] [Code]

9. **Howard Yen**, Tianyu Gao, Jinhyuk Lee, and Danqi Chen. "MoQA: Benchmarking Multi-Type Open-Domain Question Answering". In *Proceedings of the 3rd Workshop on Dialogue and Conversational Question Answering (DialDoc) @ ACL*, 2023. [Paper] [Code]

* denotes equal contribution.

Industry Experience

Meta Reality Labs — Seattle, Washington
Software Engineering Intern — Summer 2022

Facebook AI Applied Research — Menlo Park, California
Software Engineering Intern — Summer 2021

Teaching

- **Graduate Teaching Assistant** at Princeton University — Spring 2024
  Natural Language Processing (COS484)
- **Graduate Teaching Assistant** at Princeton University — Fall 2023
  Introduction to Machine Learning (COS324)
- **Research Instructor** at Princeton University — Summer 2023
  Princeton AI4ALL Summer Camp
- **Undergraduate Course Assistant** at Princeton University — Spring 2022, Spring 2023
  Natural Language Processing (COS484)
- **Undergraduate Course Assistant** at Princeton University — Spring 2020 – Fall 2022
  Algorithms and Data Structures (COS226)

Honors and Awards

- The William A. Dippel '50 *55 Graduate Fellowship — 2025
- Tau Beta Pi — 2023
- Sigma Xi — 2023
- Sigma Xi Book Award — 2023
- Phi Beta Kappa — 2022–2023
- Outstanding Student Teaching Award — 2023
- International Collegiate Programming Contest (ICPC) North America Finalist — 2021
- Shapiro Prize for Academic Excellence — 2021
- Citadel Terminal Live 2nd Place — 2020
- North Dallas Toyota Scholarship — 2019

Invited Talk

- NVIDIA, "HELMET: How to Evaluate Long-Context Language Models Effectively and Thoroughly" — 2024/10
- Sierra.AI, "Enabling Large Language Models to Generate Text with Citations" — 2023/08

Professional Services

Reviewer

- International Conference on Learning Representations (ICLR): 2025.
- ACL Rolling Review (ARR): February 2025.
- ICML 2024 Workshop on Long Context Foundation Models (ICFM).